

Introduction to C++

Basic program building blocks, control structures, functions

Marco Frailis

INAF – Osservatorio Astronomico di Trieste

The C++ language 1/2

- C++ is designed to be a statically typed, general purpose language, to be as compatible with the C language as possible
- It is designed to support multiple programming styles:
 - Procedural programming
 - Data abstraction
 - Object-oriented programming (OOP)
 - Generic programming (with templates)
 - Functional programming (to some extent, e.g. lambda functions)
- C++ is a compiled language
 - Source text files processed by a compiler producing object files
 - Object files combined by a linker yielding an executable program
 - The executable program is created for a specific hardware/system combination

The C++ language 2/2

- It provides a standard library, extending the core language, including:
 - Input/output classes
 - Data structures: dynamic arrays, linked lists, binary trees, hash tables
 - Several algorithms: searching, sorting, counting, manipulating
- Memory management
 - Low level: pointers and raw arrays, allocation, deallocation
 - Higher level: allocation and smart pointers

Object-oriented programming

- An **object** packages both data and procedures that operate on the data
- Such procedures are called **member functions** or methods or operations
- The implementation of an object is specified in its **class**, which defines the internal data (or **attributes**) of the object and the **operations** that the object can perform
- New classes can be defined as **subclasses** of a **parent class**, **inheriting** its attributes and operations
- An **abstract class** is one whose main purpose is to define a **common interface** for its subclasses
- When inheriting from an abstract class, we speak about **polymorphism**, since the same interface is associated to different implementations through its subclasses.

Generic programming

- Generic programming parameterizes algorithms so that they work for a **variety** of **types** and **data structures**
- For instance, data structures holding a collection of elements of some kind, such as vectors, lists, queues, associative containers, are general concepts and should be independent of the type of elements
- Also, algorithms to sort, copy or search a sequence of elements should be independent of the particular type of the elements or type of the container
- They can be defined as **templates**, parameterized by the types to which type are applied
- In C++, templates are a **compile time** mechanism, avoiding any run-time overhead

C++ standard evolution 1/2

- First ISO standard in 1998 (C++98)
- New revision of the standard in 2003 (C++03)
 - No new language features, just a *bug fix* release for compiler writers
- C++ Technical Report 1 (TR1) in 2005
 - A document proposing additions to the standard library

C++ standard evolution 2/2

- New C++ ISO standard in 2011 (C++11)
 - Additions to the core language, including:
 - Deducing the type of an object from its initializer (**auto**)
 - Lambda expressions
 - Move semantics
 - The range-for statement
 - Type aliases
 - Additions to the standard library, including
 - Hashed containers
 - Basic concurrency library (threads)
 - Regular expression library
 - tuple library
 - Unique and shared pointers
- Last C++ standard revision in 2014 (C++14)
 - Minor revision with small improvements, e.g. function return type deduction

C++ compilers (with C++11 conformance)

- GCC (GNU Compiler Collection), version 4.8 or greater
 - Linux, Mac OS X, Windows (with MinGW or Cygwin), iOS, Android
 - ISO support status:
<https://gcc.gnu.org/onlinedocs/libstdc+/manual/status.html#status.iso.2014>
- Clang, version 3.3 or greater
 - Linux, Mac OS X, Windows (with MinGW or Cygwin), iOS, Android
- Intel(R) C++ Compiler (commercial), version 13.0 or greater
 - Linux, Mac OS X, Windows, Android
- Microsoft Visual C++, version in Visual Studio 2013 or greater
 - Free version (with license limitations): Visual Studio Community 2013
 - Windows, Android, iOS

C++ in Astronomy

- Planck
 - Most of the processing pipelines developed in C++ (e.g. telemetry processing, calibration)
- LSST
 - Core processing developed in C++
- ALMA
 - ACS software: C++ implementation for the control system
- MAGIC telescopes
 - MARS, the analysis and reconstruction software is based on ROOT and written in C++
- Euclid
 - C++ and Python selected for developing all the processing levels
- FITS standard
 - cfitsio (C)
 - CCfits (C++)

References

● On-line references and FAQ:

- <http://www.cplusplus.com>
- <http://en.cppreference.com/w/>
- <https://isocpp.org/faq>
- <http://www.stroustrup.com/C++11FAQ.html>

● On-line tutorials

- <http://www.tutorialspoint.com/cplusplus/index.htm>
- <http://www.cplusplus.com/doc/tutorial/>
- <http://www.learncpp.com/>
- <http://www.cprogramming.com/tutorial/c++-tutorial.html>

● Books (in suggested order)


- Accelerated C++: Practical Programming by Example, by A. Koenig and B. E. Moo
- The C++ Standard Library: A Tutorial and Reference, 2nd edition, by N. M. Josuttis
- C++ Primer (5th edition), by S. B. Lippman, J. Lajoie and B. E. Moo
- The C++ Programming Language, 4th edition, B. Stroustrup

● More advanced books

- Effective Modern C++: 42 Specific Ways to Improve Your Use of C++11 and C++14, S. Meyers
- C++ Templates: The Complete Guide, by D. Vandevoorde and N. M. Josuttis
- Exceptional C++: 47 Engineering Puzzles, Programming Problems, and Solutions, by E. Sutter
- More Exceptional C++: 40 New Engineering Puzzles, Programming Problems, and Solutions, by E. Sutter

Integrated Development Environments (IDEs)

- Code::Blocks (<http://www.codeblocks.org/>)
 - A lightweight IDE (it is provided in the virtual machine created for this course)
- Eclipse CDT (<https://eclipse.org/cdt/>)
- NetBeans IDE (<https://netbeans.org/features/cpp/>)
- Emacs (<http://www.gnu.org/software/emacs/>)
- Coding Ground (http://www.tutorialspoint.com/compile_cpp11_online.php)
 - An on-line web IDE, for short tests and examples
- Others:
 - QT Creator, Xcode (Mac OS X), Visual Studio (Windows)



C++ program: main function example

ex01

```
#include <cmath>
#include <iostream>
```

Including standard library functions and classes

```
using namespace std;
```

```
const double PI = 3.1415926535897931;
```

Constants definitions

```
const double RAD2DEG = 180.0/PI;
```

```
// Convert Cartesian coordinates to spherical coordinates
```

```
int main()
```

Main function definition for the executable

```
{
```

```
    // Coordinates to be converted
```

```
    double x=10.3, y=-5.2, z=36;
```

```
    // Norm
```

```
    double r = sqrt(x*x + y*y + z*z);
```

```
    // Longitude
```

```
    double phi = atan2(y, x);
```

```
    phi = phi < 0 ? phi + 2*PI : phi;
```

```
    // Colatitude
```

```
    double theta = acos(z/r);
```

```
    cout << "Longitude = " << phi*RAD2DEG << " deg, Colatitude = "
         << theta*RAD2DEG << " deg, Radius = " << r << endl;
```

statements for coordinates conversion

Printing output

```
    return 0;
```

```
}
```

Compiling with GNU g++

- C++ source code should be given one of the valid C++ file extensions: .cpp, .cc, .cxx or .C, while the .c extension is reserved to C programs
- Saving the previous source code in file “coord.cpp”, in order to compile it and create an executable, we can use the GNU g++ compiler:

A terminal window with a dark title bar and standard window controls (close, minimize, maximize). The command prompt shows the command to compile coord.cpp into an executable named coord.

```
$ g++ -std=c++11 coord.cpp -o coord
```

- This command compiles the source code to machine code, saving it in the executable file “coord”
- The option “-std=c++11” enables compiler support for the 2011 ISO C++ standard

The **#include** preprocessor¹ directive

- Many C++ fundamental facilities, such as I/O, are not part of the core language
- They are part of the standard library. These facilities are requested using the **#include** directive
- Example: requesting the I/O standard library facilities and standard strings

```
#include <iostream>  
#include <string>
```

where `iostream` and `string` are called a standard headers

- Including a header file produces the same results as copying the header file into each source file that needs it
- Normally the include directives are inserted at the beginning of a source file

¹The preprocessor is a macro processor, called by the compiler to transform the program before compilation

Variable definition and initialization

- A variable (also called object in C++) provides a named storage (a region of memory) that our program can manipulate
- Each variable definition starts with a type specifier, followed by a comma separated list of one or more names (identifiers)

```
int day, month, year;  
double salary;
```

- A definition may also provide an initial value for the object (initialization)
- C++ supports four forms of variable initialization: copy-initialization, direct-initialization, list-copy-initialization, list-direct-initialization

```
int month = 9, year = 2012, day;           // copy-initialization  
double salary(1250.23);                    // direct-initialization
```

OK for built-in types

```
C++11 std::complex<float> c = {-1.0, 0.0}; // list-copy-initialization  
long int count{35600};                    // list-direct-initialization
```

Recommended for class types

- It is safer to initialize every object of built-in type

```
// Boolean type
bool empty = true;

// Character types
char single_letter = 'D';

// Integer types
char tiny_value = 127;
short small_value = 32767;
int value = 2147483647;
long large_value = 2147483647;
long long huge_value = -9.22E18LL;

// Unsigned integer types
unsigned char r = 255;
unsigned short small_pvalue = 65535;
// ...etc.

// Floating point types
float temperature = 2.323787;
double theta = 2.663832728147556;
long double ltheta = 2.6638327281475567373L;
```

- Additional character types: `wchar_t` and `char16_t` and `char32_t`, for Unicode characters (UTF-16 and UTF-32)

Literals

- Literal integer constants

- Notations:

```
242 /* decimal */    0362 /* octal */    0xF2 /* hexadecimal */
```

- Integer literal types: int or long by default, depending on value.
Use of suffix to force type

```
128U /* unsigned int */    12ULL /* unsigned long long */
```

- Floating point literals: double by default

```
2.323787F /* float */    2.6638327281475567373L /* long double */
```

- Character literals (char type, ASCII): 'a' '3'

- Nonprintable characters: use of escape sequence

```
'\n' /* newline */    '\t' /* horizontal tab */    '\?' /* question mark */  
'\'' /* single quote */    '\"' /* double quote */    '\0' /* null character */
```

- String literals are *arrays* of characters

```
"\tHello World\n"    "" /* empty string */
```

Expression: Arithmetic, relational and logical operators

- An expression is composed of one or more operands and (usually) an operator; it produces a result

- Arithmetic operators:

$a + b$
 $a * b$

$a - b$
 a / b
 $a \% b$

Reminder after
division a / b

- Relational operators:

$a < b$
 $a > b$
 $a == b$

$a <= b$
 $a >= b$
 $a != b$

- Logical operators

Operator	Meaning	Example
<code>&&</code>	and	<code>(a >= 0) && (a < 5)</code>
<code> </code>	or	<code>(a < 0) (a >= 5)</code>
<code>!</code>	not	<code>!((a >= 0) && (a < 5))</code>

Bitwise operators

- Bitwise operators:

Operator	Meaning	Example
<<	left shift	a << 3
>>	right shift	a >> 2
&	bitwise and	a & mask
	bitwise or	a mask
^	bitwise xor	a ^ mask
~	bitwise not	~mask

```
unsigned char a = 20, mask = 1; // i.e. a = 00010100, mask = 00000001
```

```
a >> 2           // result = 00000101
mask << 3         // result = 00001000
~mask            // result = 11111110
```

```
a & (mask << 4) // result = 00010000 (check if bit 4 is set)
a | (mask << 1) // result = 00010110 (set bit 1)
a ^ (mask << 3) // result = 00011100 (flip bit 3)
```

Assignment and compound assignment operators

```
// definition and initialization
int index = 0 , step = 5;
double deltaT = 0;

// assignment
deltaT = 0.3;
index = step * deltaT ; // here an implicit truncation is performed
```

- The result of an assignment is the left-hand operand
- We often apply an operator to an object and reassign the result to the same object:

```
index = index + step * deltaT;
```

- C++ provides compound-assignment operators as a shorthand, for arithmetic and bitwise operators

```
index += step * deltaT; // equivalent to the previous statement
```

Increment and decrement operators

- Increment (++) and decrement (--) operators can be used as a shorthand for adding or subtracting 1 from an object
- There are two forms of these operators: prefix and postfix
- The prefix form increment (decrement) its operand and yields the changed value as its result
- The postfix form increment (decrement) its operand but yields a copy of the original, unchanged value as its result

```
int i = 0, j;
```

```
j = ++i;    // j = 1 , i = 1;  
j = i++;    // j = 1 , i = 2;
```

Implicit type conversions

- In expressions with operands of mixed types, the types are converted to a common type by the compiler
- For arithmetic conversions, the rules define a hierarchy of conversions in which operands are converted to the widest type in the expression
- The conversion rules are defined so as to preserve the precision of the values involved
- Conversion to/from **bool**:
 - Integral and floating-point values can be converted to **bool**. If the value is 0, the resulting bool is **false**, otherwise it is **true**.
 - **bool** can be converted to other types: **true** is converted to 1, **false** to 0

Explicit conversions

- An explicit conversion is spoken of as a **cast**

- C old-style cast:

```
int x = 37, y = 6;  
float result = (float) x / y;    // or float(x) / y
```

- Casts should be discouraged: they turn off normal type-checking
- C++ tries to make casts more visible providing named cast operators:

```
result = static_cast<float>(x) / y;
```

- Additional C++ named cast operators:

dynamic_cast, const_cast, reinterpret_cast

Strings in C++

- The Standard Library provides the `string` type to support variable length character strings
- The library takes care of managing the memory associated with storing the characters
- It provides several ways to initialize a `string` variable:

```
#include <string>
```

```
int main()
{
    std::string s1;           // s1 is an empty string
    std::string s2(s1);       // s2 is initialized as a copy of s1
    std::string s3("swing");  // s3 is a copy of the string literal
    std::string s4(5, '-');   // s4 = "-----"
}
```

- `string` is part of the `std` namespace
- A namespace is a collection of related names: the standard library uses `std` to contain all the names that it defines
- `std::string` is a qualified name, using the scope operator (`::`)

String operations

```
#include <string>

using namespace std;

int main()
{
    string s1;
    string s2 = "Hello";
    string s3 = s2 + " World";           // string + literal concatenation

    bool isEmpty = s1.empty();           // check if s1 is an empty string
    char first = s3[0];                   // get the first character
    s3[4] = 'o';                           // modify the 5th character in s3
    bool isLess = s2 < s1;                 // comparison operator between strings
                                         // (lexicographic order)
    s1 = s2;                               // copy s2 in s1, replacing s1 characters
    string s4 = s3.substr(6, 5);           // get the substring starting at index 6
                                         // and taking 5 chars (s4 = "World")
    string::size_type size = s4.size();    // get string size in bytes
}
```

Type defined in the string class:
synonym of an unsigned int, but more portable

C++ Input/Output

- The standard library defines a family of types that support I/O to and from devices such as files and console windows. Additional types allow strings to act like I/O streams
- The I/O types are defined in three separate **headers**:
 - **<iostream>** defines the types used to read from and write to a console window
 - **<fstream>** defines the types used to read from and write to named files
 - **<sstream>** defines the types used to read from and write to in-memory strings
- The **iostream** header includes the definition of three I/O **objects**:
 - The **istream** object named **cin**, also referred to as *standard input*
 - The **ostream** object named **cout**, also referred to as *standard output*
 - The **ostream** object named **cerr**, also referred to as *standard error*

Standard I/O example

```
#include <iostream>
using namespace std;

int main()
{
    cout << "Please, insert two numbers:" << endl;
    float v1, v2;

    cin >> v1 >> v2;

    cout << "The sum of " << v1 << " and " << v2
         << " is " << v1 + v2 << endl;
    return 0;
}
```

- The stream operators (<< and >>) are left associative and return the left operand, i.e. the stream object)
- So we can chain the stream operations

Formatting floating point values

```
#include <iostream>
#include <iomanip>
using namespace std;
```

I/O format state (parametric)
manipulators

```
int main()
{
    const double PI = 3.1415926535897931;

    cout << PI << endl;
    cout << setprecision(15) << PI << endl;
    cout << scientific << PI << endl;
    cout << fixed << setw(20) << PI << endl;    // setw sets the field width
    cout << resetiosflags(ostream::floatfield); // revert to default notation

    cout << PI << endl;
}
```

Output:

```
3.14159
3.14159265358979
3.141592653589793e+00
  3.141592653589793
3.14159265358979
```

Statements

- Most statements in C++ end with a semicolon
- We have already seen some expression statements
 - i.e. an expression followed by a semicolon
 - Commonly expression statements affect the program's state: assignment, increment, input/output operators, declaration statements
- A compound statement, or block, is a (possibly empty) sequence of statements surrounded by a pair of curly braces
- Compound statements can be used where the rules of the language require a single statement

Scope of a name

- Every name in a C++ program must refer to a unique entity (e.g. a variable, function, type, etc.)
- A name can be reused, also with different meanings, as long as it is used in different context
- Such context is the **scope**, i.e. a region of the program
- There are different kinds of scope:
 - The statements between a pair of matching braces form a scope. The body of the `main` or the body of every function and the scopes nested inside a function (such as a **block**) form **local scopes**
 - Names defined outside any function have **global scope**. They are accessible from anywhere in the program
 - **Namespaces** partition the global namespace. A namespace is a scope, as we have already seen with the standard library namespace `std`
 - Every **class** defines its own new scope. An examples is given by the type `size_type` of the `string` class

Scope example

```
#include <iostream>
```

```
using namespace std;
```

```
const string s = "-----";
```

Global const variable

```
int main()
{
    unsigned int s = 0;
```

s is a variable in the main function local scope

```
{
    unsigned int s = 1;
    cout << "Inner s: " << s << endl;
}
```

In this block, s definition hides the outer one

```
cout << "Outer s: " << s << endl;
cout << "Global s: " << ::s << endl;
```

```
return 0;
```

```
}
```

● Output:

```
Inner s: 1
Outer s: 0
Global s: -----
```

if statement

Condition

```
if (rawVal > 90) {  
    cout << "Warning: hard limit exceeded" << endl;  
    ++hardLimit ;  
}  
else if (rawVal > 70)  
    ++softLimit;  
else {  
    sum += rawVal;  
    ++numSamples;  
}
```

Statement 1

Statement 2

- Second **if** form: without the **else** branch

```
if (i < j)  
    ++i;
```


while and do while statements

while statement

Condition

```
while (i > 0 && v[i] < key) {  
    v[i+1] = v[i];  
    --i;  
}
```

Statement

do-while statement

```
double threshold = -1;  
do {  
    cout << "Please enter a threshold level in the range [0, 1]"  
        << endl;  
    cin >> threshold;  
} while (threshold < 0 || threshold > 1);
```

Statement executed
at least once

```
int i = 0;  
while (i < size) {  
    // operations that don't  
    // change the value of i  
    ++i;  
}
```

Recurrent pattern:
for loop

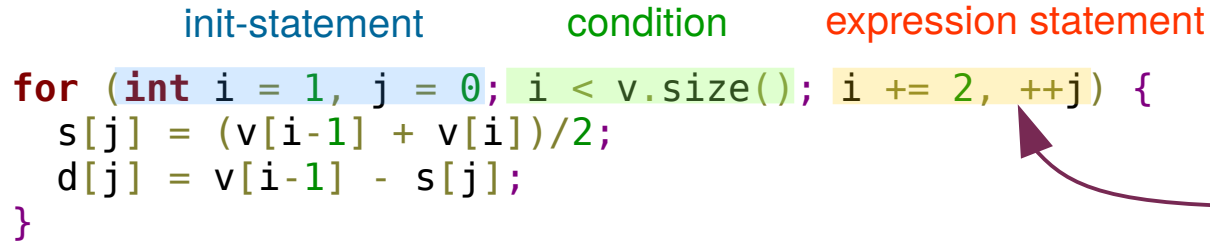


for and range-based for loops

for statement

init-statement condition expression statement

```
for (int i = 1, j = 0; i < v.size(); i += 2, ++j) {  
    s[j] = (v[i-1] + v[i])/2;  
    d[j] = v[i-1] - s[j];  
}
```

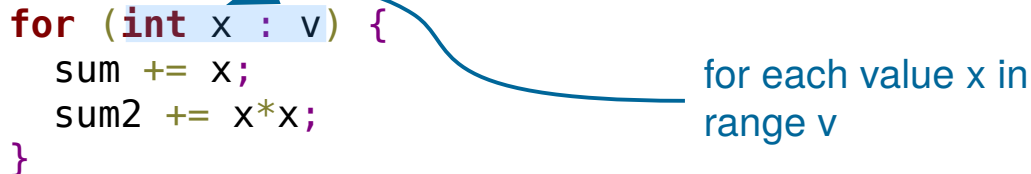


Executed after body
of for loop

range-for statement

```
for (int x : v) {  
    sum += x;  
    sum2 += x*x;  
}
```

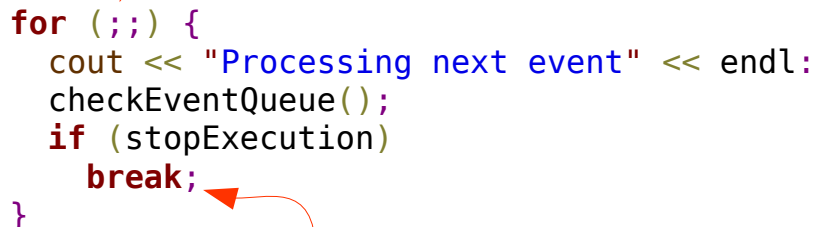
for each value x in
range v



C++11

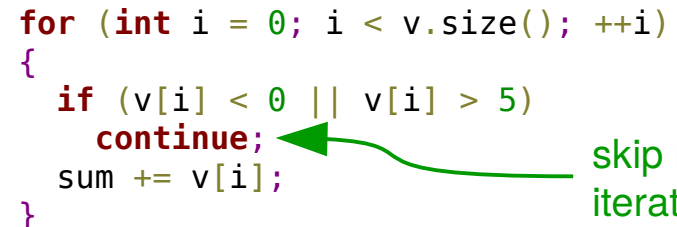
Infinite loop until event

```
for (;;) {  
    cout << "Processing next event" << endl;  
    checkEventQueue();  
    if (stopExecution)  
        break;  
}
```



ends the nearest enclosing loop


```
for (int i = 0; i < v.size(); ++i)  
{  
    if (v[i] < 0 || v[i] > 5)  
        continue;  
    sum += v[i];  
}
```



skip rest of
iteration

switch-case statement

```
int apId = packet.apId();  
switch (apId) {  
    case 1536:  
    case 1538:  
        ++lfiHkCounter; // LFI housekeeping telemetry found  
        break;  
    case 1540:  
        ++lfiSciCounter; // LFI scientific telemetry found  
        break;  
    default:  
        cout << "Unknown telemetry packet" << endl;  
}
```

A red arrow points from the text "expression with integral result" to the variable 'apId' inside the parentheses of the switch statement. The entire code block is enclosed in a dashed yellow border.

- It provides a more convenient way to write deeply nested if/else logic
- The result of the expression is compared with the value associated with each **case**
- Execution starts with the first statement following the matching label, till a **break** is found

Standard vector 1/3

- The standard library provides a type, named vector, that holds a sequence of values of a given type and grows as needed
- It is defined using a language feature called template classes

```
#include <iostream>
#include <vector>
#include <complex>
```

```
using namespace std;
```

```
int main()
{
```

Template
parameter



```
    int n = 10;
```

```
    vector<float> a(n, 5.0); // a has n elements that are copies of 5.0
    vector<int> b = {3,2,7,11,23}; // b contains the elements provided in the list
```

```
    cout << "First and last elements of b: "
         << b[0] << " " << b[b.size()-1] << endl;
```

```
    // vector from a list of numbers gathered from the standard input
```

```
    double value;
```

```
    vector<double> c;
```

```
    cout << "Please insert some values (use Ctrl-D to end): " << endl;
```

```
    // The user ends inputting numbers with the EOF character
```

```
    while (cin >> value)
```

```
        c.push_back(value); // appends a new element at the end of vector c
```

Continue ...

C++11

Standard vector 2/3



```
// Print vector elements
cout << "c elements: ";
for (double x : c)
    cout << x << " ";
cout << endl;

vector<double> d(c);    // defines d as a copy of c

// e holds a copy of the first half of d
vector<double> e(d.begin(), d.begin() + d.size()/2);

c.clear();    // removes all elements in c; c.size() == 0

// erase elements of first half of d
d.erase(d.begin(), d.begin() + d.size()/2);


c = e;    // c is now a copy of e

// prepend elements of a in e
e.insert(e.begin(), a.begin(), a.end());

vector<complex<double>> f;
f.push_back({1, -1});    // analogous to push_back(complex<double>(1,-1)), but shorter
                          // thanks to initializer list
C++11 f.emplace_back(1, -1);    // emplace_back is analogous to push_back but it just needs the
                                // arguments to construct the value, which is created in-place
}
```

- vector is one of the sequential containers provided by the standard library
- To support fast random access to the elements, vector elements are stored contiguously
- Vectors grow dynamically: they allocate capacity beyond what is immediately needed, as a reserve for new elements
 - When the capacity is exceeded, a new block of contiguous memory is automatically allocated (e.g. by a factor of 2 larger than the previous one)
 - Vector class provides two member functions concerning its memory block: `vector::capacity()` and `vector::reserve(n)`

```
vector<int> a;  
for (int i = 0; i < 33; ++i) {  
    a.push_back(i);  
    cout << "a capacity: "  
        << a.capacity()  
        << endl;  
}
```



`a.reserve(n)` ← lets us set the initial vector capacity

The **auto** type specifier and type aliases

- When a definition of a variable has an initializer, we can let the variable have the type of the initializer using the **auto** type “placeholder”

```
vector<complex<double>> v = {{1,0}, {1,-1}, {0,1}, {0,-1}};
```

C++11

```
auto size = v.size();
```

The **size** variable has type
`vector<complex<double>>::size_type`

```
cout << "v elements: ";
```

```
for (auto x : v)
```

```
    cout << x << " ";
```

```
cout << endl;
```

x has type
`complex<double>`

- When defining a variable with **auto**, prefer the “=” syntax
- Sometime we need a new name for a type, e.g. because it is too long or complex or because we need different types to have same name

```
using vec_compd = vector<complex<double>>;
```

C++11

```
vec_compd v = {{1,0}, {1,-1}, {0,1}, {0,-1}};
```

Older syntax using **typedef**

```
typedef vector<complex<double>> vec_compd;
```

Functions

- Functions are named computations. We can identify important parts of our problems and create named computations corresponding to those parts
- A function is defined by specifying its return type, followed by the function name and then by a parameter list enclosed in () and finally the function body, which is enclosed in {}
- A function is uniquely represented by its name and the list of parameter types. They form the so called function signature

```
int gcd(int v1, int v2) // return the greatest common divisor
{
    while (v2) {
        int temp = v2;
        v2 = v1 % v2;
        v1 = temp;
    }
    return v1;
}
```


Argument passing

- In the previous example, parameters are **passed by value**, i.e. they are initialized by copying the corresponding argument
- When we pass built-in types by value, the cost of copying them into the function parameters is negligible
- However, if we need to pass a long vector or string, copying them can be time consuming
- Often, we also need to define functions that modify the arguments passed, or we need them to return additional information to the caller

References and const references

- A **reference** to an object is another name, or synonym for that object

```
vector<double> samples;
```

```
vector<double>& s = samples; // s is a synonym for samples
```

- In the example above, variable `s` is defined as a reference (synonym) to `samples`
- Any operation performed on `s` is equivalent to doing the same to `samples`, and vice versa
- A reference always refers to the object to which it was initialized (no null references)
- The main use of references is for specifying arguments and return values for functions in general
- When we add a **const** to a reference, we can only use the reference to read values from the original object, but we cannot change its values (no write access)

Argument passing: const references

```
double mean(const vector<double>& samples)
{
    auto size = samples.size();

    if (size == 0)
        throw domain_error("mean of an empty vector");

    double sum = 0;
    for (const auto& x : samples)
        sum += x;
    return sum/size;
}
```

samples is a synonym for the actual argument (a vector) passed in the function call

each element in samples is passed to x as const reference

- The type that we specify for the function argument is called “reference to **const** vector of **double**”
 - **No copy** of the actual argument passed with the function call (efficient)
 - The actual argument is passed as a read-only object
- In the range-for within the function body, x is a const reference to each element in samples (**no copy performed**)

- There are situations where passing parameters by value or using const references don't work. For instance, in a simple function that swaps two values we need to use non-const references:

```
void swap(int& v1, int& v2)
{
    int tmp = v2;
    v2 = v1;
    v1 = tmp;
}

int main()
{
    int i = 5, j = 20;
    swap(i,j);    // After this call: i = 20, j = 5
}
```

- This function has no return value. Its return type is **void**, a C/C++ built-in type used in a few restricted ways, e.g. to name a return type
- The purpose of the swap function is to change its argument values. So its parameters are declared as references

Argument passing: summary

- When function parameters are built-in types, they can be passed by value if the function does not need to change their value
- Whenever the function does not need to change the parameter's value, and the parameter has a type that can be time consuming to copy, then the parameter should be a const reference
- When the function intends to change the arguments values, the corresponding parameters must be declared as non-const references
 - Arguments passed as non-const reference parameters must be non-temporary objects (lvalues)